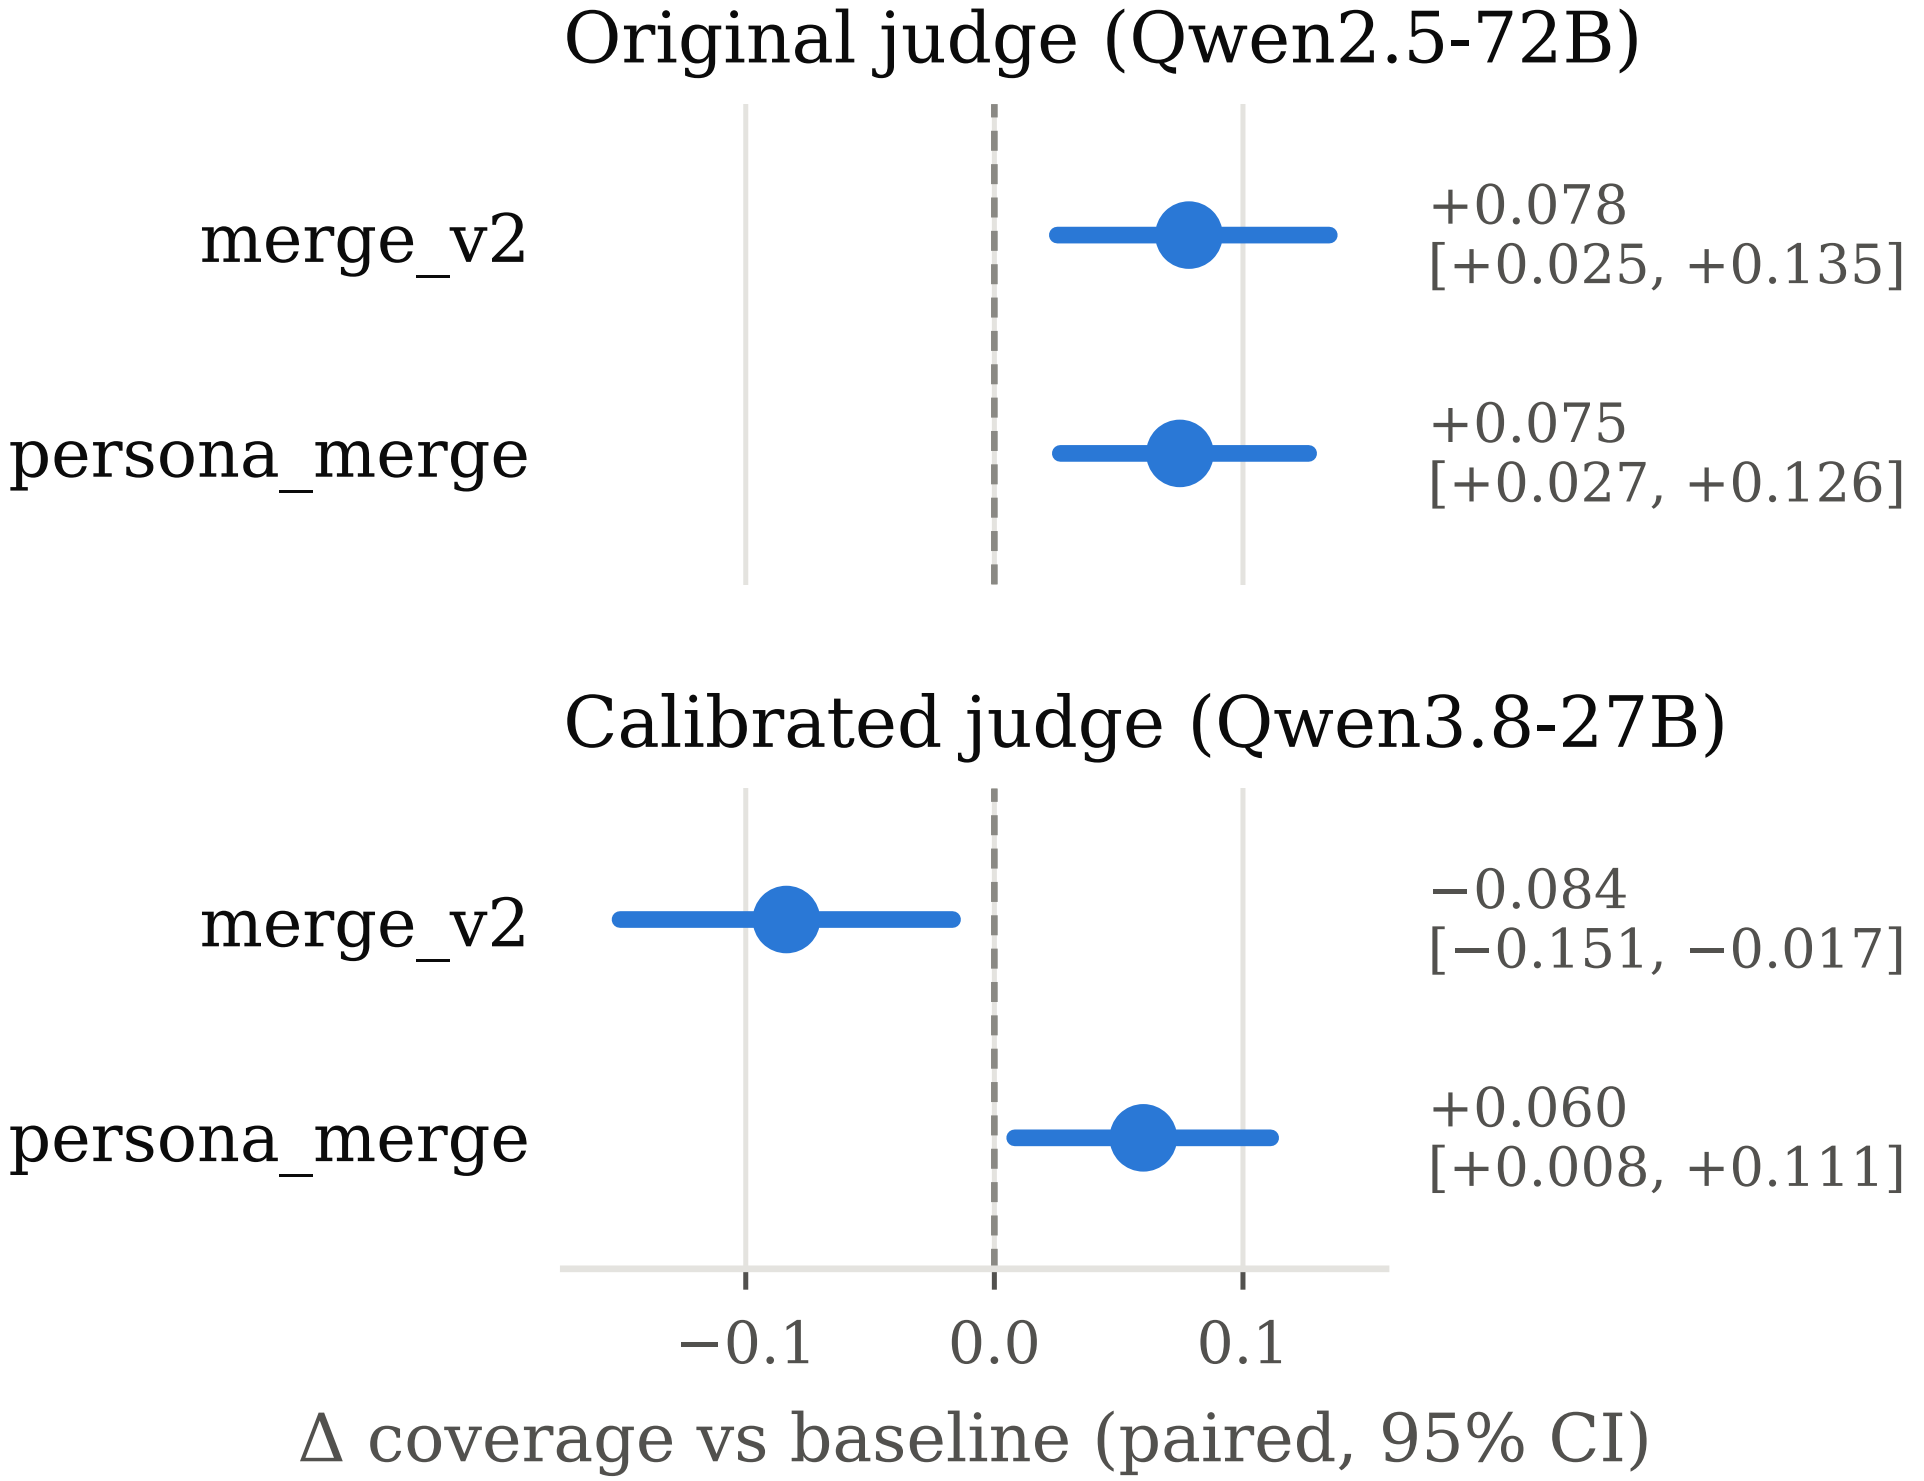
Original judge (Qwen2.5-72B)
merge_v2
+0.078
[+0.025, +0.135]
persona_merge
+0.075
[+0.027, +0.126]
Calibrated judge (Qwen3.8-27B)
merge_v2
−0.084
[−0.151, −0.017]
persona_merge
+0.060
[+0.008, +0.111]
−0.1
0.0
0.1
Δ coverage vs baseline (paired, 95% CI)